

TASK 1: Oracle VirtualBox Setup and Ubuntu

Date: 12-02-2026

Oracle VM VirtualBox is a virtualization application that enables users to run different operating systems on one computer. This involved first installing Ubuntu Linux inside the VirtualBox to create a virtual Linux environment. It is used for:

- Operating system and command line operation
- Create a secure and isolated environment for testing purposes

1. Installing Oracle VirtualBox

Steps:

- 1: Visit <https://www.virtualbox.org>
- 2: Download the installer for Windows hosts.
- 3: Open the installer file and click Next to start the setup process.
- 4: Keep the default configuration settings and install all the required components.
- 5: Allow the network driver installation.
- 6: Hit Finish to finalize the installation.

2. Download Ubuntu

Steps:

- 1: Go to <https://ubuntu.com/download/desktop>
- 2: Download the Ubuntu 22.04 LTS ISO file.

3. Creating Ubuntu Virtual Machine

Steps:

- 1: Open VirtualBox and click on New.
- 2: Enter the following information:
 - Name: Ubuntu
 - Linux type
 - Version: Ubuntu (64-bit)
- 3: Allocate memory (RAM):
 - Minimum: 2048 MB
 - Recommended: 4096 MB
- 4: Set up a virtual hard disk:
 - Disk Type: VDI (VirtualBox Disk Image)
 - Storage Allocation: Dynamically allocated
 - Disk Size: Between 20–25 GB

4. Installing Ubuntu

Steps:

- 1: Select the created virtual machine and click on Start.
- 2: Select the ISO with the download Ubuntu operating system.
- 3: Click Install Ubuntu.
- 4: Select the Normal installation option.
- 5: Select Erase disk and install Ubuntu (no changes will be made on your physical disk; alterations will be limited to the virtual disk only).
- 6: Provide a username and password.
- 7: When installation is done, the user should restart the virtual machine.




Figure 1: Ubuntu OS installed in a VM




Figure 2: USER PROFILE




Figure3: USER TERMINAL

Task 2 : Working Docker Images and NGC Container

Date: 18-02-2026

NVIDIA NGC Login and API Key Generation

To use GPU-enabled containers, which are available on the NVIDIA Container Registry, an **NVIDIA NGC** account was created. To ensure that the system can securely download the NGC containers, an API key has been generated and used to authenticate Docker.

Step 1: Create NVIDIA NGC Account

1. Launch a web browser.
2. Visit <https://ngc.nvidia.com>
3. Press **Sign In**.
4. Click on Create Account.
5. Fill out the registration form, and verify your email to activate your account.

Step 2: Generate NGC API Key

1. Log in to the NVIDIA NGC dashboard.
2. Click your User Profile icon.
3. Select **API Key** option.
4. Click Generate API Key.
5. Copy the generated key and save it somewhere secure for later use.

Step 3: Login to NGC Using Docker (Ubuntu Terminal)

In the Ubuntu terminal, run the following command: Code: docker login nvcr.io Input the following credentials when asked

- **Username:** \$oauthtoken
- **Password:** Paste the generated NGC API key

With a successful login process, the terminal will display: The terminal after a login sequence is as follows: Login Succeeded.

This shows that the system is successfully authenticated and can now pull containers from the NVIDIA NGC registry.

NGC Catalog

All you need to build AI—GPU-optimized containers, pretrained models, SDKs, and Helm charts—unified in one catalog for cloud, data-center, or edge.

Top Technology

The most-downloaded AI frameworks, models, and SDKs on NGC—curated by NVIDIA, battle-tested by the community, ready to accelerate your next project.

[All](#)
[Containers](#)
[Collections](#)
[Models](#)
[Resources](#)
[Helm Charts](#)

[View More](#)

NVIDIA PyTorch

PyTorch is a GPU accelerated tensor computational framework. Functionality can be extended with common...

High Performance Comput...

Natural Language Understand...

 Container

Updated 22 days ago

NVIDIA Triton Inference Server

Triton Inference Server is an open source software that lets teams deploy trained AI models from any framework, from...

Object Detection

Automatic Speech Recognition

 Container

Updated 20 days ago

NVIDIA cuda

Container registry for CUDA images

CUDA

 Container

Updated a month ago

NVIDIA DeepSeek-R1

NVIDIA NIM for GPU accelerated DeepSeek-R1 inference through OpenAI compatible APIs

 +1

Collection

NVIDIA Cosmos World Foundation Models

Cosmos World Foundation Models: A family of highly performant pre-trained world foundation models purpose...

 +1

Collection

NVIDIA DeepStream SDK

NVIDIA DeepStream SDK enables developers to build accelerated pipelines for a wide range of use-cases such as...

Figure 4: NVIDIA NGC USER DASHBOARD

- Dashboard
- Audit
- Organization Profile
- Teams
- Usage
- Activate Subscription
- Subscriptions
- Webhook Management

Setup > API Keys

API Keys

[+ Generate Personal Key](#)

Personal Keys

 Search table

1 Result [View](#) ▾

Key ID	Key Name	Expiration	Key Value	Status	
87357941-2298-4b21-ab36-356dde07fcac	AFEF AHMED SHAIK	02/25/2027	nvapi-*****g5q	ACTIVE	⋮

25 ▾ Items

<< < 1 > >>

Go to 1

NGC API Keys

NVIDIA NGC API keys are required to authenticate to NGC services using NCG CLI, Docker CLI, or API communication. NVIDIA NGC supports two types of API keys.

Personal API Key

Any user who is a member of an NGC org can generate Personal API Keys. These keys are tied to the user's lifecycle within the NGC org and can access up to the permissions and services assigned to the user. During the key generation steps, users can configure which NGC services are accessible by the API key and the time-to-live from one hour to 'never expires'.

Legacy API Key

This is the original type of API key available in NGC since its inception. This type allows you to create only one 'API key' at a time. Generating a new key automatically revokes the previous one, as they cannot be rotated. The active key immediately becomes invalid when you create a new key.

Figure 5: NVIDIA NGC API ORGANIZATION DASHBOARD

Organization

Dashboard

Audit

Organization Profile

Teams

Usage

Activate Subscription

Subscriptions

Webhook Management

Setup > API Keys

API Keys

Generate Personal Key

Personal Keys

Search table

Key ID	Status
87357941-2298-4b21-ab3	ACTIVE
be3c991a-625d-4906-9cd	ACTIVE

25 Items

NGC API Keys

NVIDIA NGC API keys are required to authenticate to NGC services using NCG CLI, Docker CLI, or API communication. NVIDIA NGC supports two types of API keys.

Personal API Key

Any user who is a member of an NGC org can generate Personal API Keys. These keys are tied to the user's lifecycle within the NGC org and can access up to the permissions and services assigned to the user. During the key generation steps, users can configure which NGC services are accessible by the API key and the time-to-live from one hour to 'never expires'.

Legacy API Key

Generate Personal Key

The following Personal API Key has been successfully generated for use with this organization. This is the only time your key will be displayed.

\$ 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

Keep your Personal Key a secret. Do not share it or store it in a place where others can see or copy it.

CloseCopy Personal Key

Figure 6: NVIDIA NGC API KEY GENERATION

Docker Images Before Pulling GitHub and Friend Images

The existing Docker images were checked first to ascertain the current status of the overall Docker trends before downloading any new images.

The below command was run:

```
docker images
```

This command lists out all the docker images stored locally in the system. These details include repository, tag, image ID, the date when the image was created, and the size of the image.

By running this command, verifying which images are already present and which have been downloaded later becomes much easier.

At this point, only previously pulled images (for example, NVIDIA container images) were in the system. Second, there were no custom images downloaded into the system created by the user or their peers.




Figure 7: Docker Images Before Pulling GitHub and Friend Images

Pulling TensorFlow Container & Program Execution

The TensorFlow container was retrieved from the NVIDIA NGC registry to initiate a pre-configured deep learning Dockerized environment. This container has Python, TensorFlow, and all the necessary libraries needed in the process, hence, there is no need to install dependencies manually.

```
dockerpullnvcr.io/nvidia/tensorflow:25.10-tf2-py3
```

After the downloading process had been completed, the container was then launched in interactive mode through:

```
dockerrun -it nvcr.io/nvidia/tensorflow:25.10-tf2-py3
```

This command started an interactive Python session inside the container.

To verify the TensorFlow installation, whether it was successful or not, and if the software behaved as required, the following commands were run in Python:

```
import tensorflow as tf  
print("TensorFlow Version:", tf.__version_)  
print("GPU Available:", tf.config.list_physical_devices('GPU'))
```

The output included the currently installed tensorflow version and a statement that the container was running properly. Since the system does not have an NVIDIA GPU, the GPU list was empty, proving that the GPU has been disabled and that TensorFlow is working in CPU mode.

This verified that the TensorFlow container was successfully downloaded and run within the Docker environment.

Ask Gordon BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit BETA

Docker Hub

Docker Scout

Extensions

Images [Give feedback](#)

LocalMy Hub

42.12 GB / 30.43 GB in use3 Images

Last refresh: 2 hours ago

Search

	Name	Tag	Image ID	Created	Size	Actions
	nvcr.io/nvidia/tensorflow	25.01-tf2-py3	4b15081766f8	1 year ago	31.97 GB	
	nvcr.io/nvidia/pytorch	26.01-py3	38ed2ecb2c16	1 month ago	30.43 GB	

Terminal

Error response from daemon: failed to resolve reference "nvcr.io/nvidia/tensorflow:26.01-tf2-py3": nvcr.io/nvidia/tensorflow:26.01-tf2-py3: not found
PS C:\Users\afeef> docker pull nvcr.io/nvidia/tensorflow:25.01-tf2-py3
25.01-tf2-py3: Pulling from nvidia/tensorflow
4f1559b21c0f: Pulling fs layer
e852726cdf4b: Pulling fs layer
a120e6e82007: Pulling fs layer
7dfa49293b99: Pulling fs layer
81355a82f4ed: Pulling fs layer
07e412f82f03: Pulling fs layer
de44b265507a: Pulling fs layer
8ec2f291c690: Pulling fs layer
2dabfd5dddb1: Pulling fs layer
4eb34e451322: Pulling fs layer
75f58769314d: Pulling fs layer
4f4fb700ef54: Pulling fs layer
5a9563dd6b7e: Pulling fs layer
4459897c0dc0: Pulling fs layer
e5a680752b35: Pulling fs layer

Engine running

RAM 7.29 GB CPU 0.00% Disk: 76.65 GB used (limit 1006.85 GB)

Terminal v4.62.0

Figure 8: TensorFlow Container Pull

PERSONAL

Ctrl+K

?

2

A

-

×

Ask Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Images

[Give feedback](#)

Local

My Hub

59.23 GB / 30.43 GB in use

7 Images

Name

Tag

↑

Image ID

Created

Size

Actions

nvcr.io/nvidia/tensorflow

25.01-tf2-py3

4b15081766f8

1 year ago

31.97 GB

Terminal

+

▼

×

```

PS C:\Users\afeef> docker run nvcr.io/nvidia/tensorflow:25.01-tf2-py3

=====
== TensorFlow ==
=====

NVIDIA Release 25.01-tf2 (build 134984172)
TensorFlow Version 2.17.0
Container image Copyright (c) 2025, NVIDIA CORPORATION & AFFILIATES. All rights reserved.
Copyright 2017-2024 The TensorFlow Authors. All rights reserved.

Various files include modifications (c) NVIDIA CORPORATION & AFFILIATES. All rights reserved.

This container image and its contents are governed by the NVIDIA Deep Learning Container License.
By pulling and using the container, you accept the terms and conditions of this license:
https://developer.nvidia.com/ngc/nvidia-deep-learning-container-license

WARNING: The NVIDIA Driver was not detected. GPU functionality will not be available.
Use the NVIDIA Container Toolkit to start this container with GPU support; see
https://docs.nvidia.com/datacenter/cloud-native/ .

NOTE: The SHMEM allocation limit is set to the default of 64MB. This may be

```

Windo...

×

Windo...

×

Windo...

×

Windo...

×

Windo...

×

Windo...

×

Windo...

×

Engine running

||

:

RAM 1.51 GB CPU 0.12% Disk: 89.15 GB used (limit 1006.85 GB)

Terminal

✓ v4.62.0

Figure 9: TensorFlow Container Run

PyTorch Container Pull & Program Execution

To have a ready deep learning environment inside Docker, the PyTorch container was downloaded from the NVIDIA NGC registry. Not only does this container require no manual configuration, but also it comes with Python, PyTorch, CUDA libraries (in case a compatible GPU is available), and all necessary dependencies.

The image was pulled using this command:

```
docker pull nvcr.io/nvidia/pytorch:26.10-py3
```

After the download was performed successfully, the container was launched in the interactive mode with:

```
docker run -it nvcr.io/nvidia/pytorch:26.10-py3
```

This command opened a Python terminal within the container environment.

To make sure that the PyTorch tool was correctly installed and working properly, the following Python commands were executed:

```
import torch  
print("PyTorch Version:", torch.__version__) print("GPU Available:",  
torch.cuda.is_available())
```

The output showed the installed PyTorch version and the availability of GPU acceleration. Since the system does not have an NVIDIA GPU, the result returned False, meaning that PyTorch was running on CPU.

This verified that the PyTorch container has been correctly downloaded and is running inside the Docker environment.

docker desktop

PERSONAL

Q Search

Ctrl+K

?

2

A

-

X

Ask Gordon BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit BETA

Docker Hub

Docker Scout

Extensions

Images

Give feedback

LocalMy Hub

59.23 GB / 30.43 GB in use7 Images

Last refresh: 1 day ago

Q pyt

	Name	Tag	Image ID	Created	Size	Actions
	nvr.io/nvidia/pytorch	26.01-py3	38ed2ecb2c16	1 month ago	30.43 GB	

Terminal

Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\afeef> docker pull nvr.io/nvidia/pytorch:26.01-py3
26.01-py3: Pulling from nvidia/pytorch
Digest: sha256:38ed2ecb2c16d10677006d73fb0a150855d6ec81db8fc66e800b5ae92741007e
Status: Image is up to date for nvr.io/nvidia/pytorch:26.01-py3
nvr.io/nvidia/pytorch:26.01-py3
PS C:\Users\afeef>

Windo... X
Windo... X
Windo... X
Windo... X
Windo... X
Windo... X

Engine runningRAM 1.49 GB CPU 0.06% Disk: 89.15 GB used (limit 1006.85 GB)> Terminal v4.62.0

Figure 10: PyTorch Pull

Ask Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Images

Local

My Hub

59.23 GB / 30.43 GB in use 7 images

Last refresh: 1 day ago

pyt

	Name	Tag	Image ID	Created	Size	Actions
☐	nvcrl.io/nvidia/pytorch	26.01-py3	38ed2ecb2c16	1 month ago	30.43 GB	

Terminal

PS C:\Users\afeef> docker run nvcrl.io/nvidia/pytorch:26.01-py3

=====

== PyTorch ==

=====

NVIDIA Release 26.01 (build 256811083)

PyTorch Version 2.10.0a0+a36e1d3

Container image Copyright (c) 2025, NVIDIA CORPORATION & AFFILIATES. All rights reserved.

Copyright (c) 2014-2024 Facebook Inc.

Copyright (c) 2011-2014 Idiap Research Institute (Ronan Collobert)

Copyright (c) 2012-2014 Deepmind Technologies (Koray Kavukcuoglu)

Copyright (c) 2011-2012 NEC Laboratories America (Koray Kavukcuoglu)

Copyright (c) 2011-2013 NYU (Clement Farabet)

Copyright (c) 2006-2010 NEC Laboratories America (Ronan Collobert, Leon Bottou, Iain Melvin, Jason Weston)

Copyright (c) 2006 Idiap Research Institute (Samy Bengio)

Copyright (c) 2001-2004 Idiap Research Institute (Ronan Collobert, Samy Bengio, Johnny Mariethoz)

Copyright (c) 2015 Google Inc.

Copyright (c) 2015 Yangqing Jia

Copyright (c) 2013-2016 The Caffe contributors

All rights reserved.

Window...

Window...

Window...

Window...

Window...

Window...

Engine running

RAM 1.54 GB CPU 0.00% Disk: 89.19 GB used (limit 1006.85 GB)

Terminal v4.62.0

Figure 11: PyTorch Run

Use Case Container Setup Using NGC Pull and Run

As a practical example to show the downloading and execution processes of NGC containers with Docker, the Triton Inference Server container was used. It is built in with accelerated libraries for efficient, high-performance inference, and optimized to deploy and serve AI models in production settings.

The container image was pulled from the NVIDIA NGC registry using the below command; This is followed by pulling the specific container image from the NVIDIA NGC registry:

```
dockerpullnvcr.io/nvidia/tritonserver:24.10-py3
```

After the image was downloaded successfully, the container was launched to confirm that it functions properly:

```
dockerrunnvcr.io/nvidia/tritonserver:24.10-py3
```

Once the container was started, initialization logs appeared on the terminal, implying that the Triton Inference Server was starting correctly. Thus, once the container started, initialization logs appeared on the terminal, and it was evident that the Triton Inference Server was starting successfully, and these logs verified that the container was functioning properly.

Triton Inference Server is common for deploying and managing trained machine learning models so that applications can send inference requests and get predictions in real time. The experiment has proven the capabilities of NGC container in real-world AI model serving and deployments.

docker.desktop

PERSONAL

Q

Search

Ctrl+K

2

A

—

×

✦ Ask Gordon BETA

📦 Containers

🔧 Images

📁 Volumes

🌐 Kubernetes

🔑 Builds

📦 Models

🔧 MCP Toolkit BETA

🌐 Docker Hub

🔍 Docker Scout

🔧 Extensions

Images [Give feedback](#)

Local

My Hub

59.23 GB / 30.43 GB in use 7 Images

Last refresh: 1 day ago

Q tri

	Name	Tag	Image ID	Created	Size	Actions
☐	nvcr.io/nvidia/tritonserver	24.01-py3	338072076104	2 years ago	22.57 GB	

Terminal

Windows PowerShell

Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! <https://aka.ms/PSWindows>

PS C:\Users\afeef> docker pull nvcr.io/nvidia/tritonserver:24.01-py3

24.01-py3: Pulling from nvidia/tritonserver

Digest: sha256:3380720761045fc16ba3bcb96cfa54034531fc302df54ecac6b2a4deeab07bbd

Status: Image is up to date for nvcr.io/nvidia/tritonserver:24.01-py3

nvcr.io/nvidia/tritonserver:24.01-py3

PS C:\Users\afeef>

Q

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Engine running

RAM 1.45 GB CPU 0.25% Disk: 89.19 GB used (limit 1006.85 GB)

>_ Terminal

✓ v4.62.0

Figure 12: Use Case Container Pull




Figure 13: Use Case Container Run

Task 3 : Installation and Pulling Packages from DOCKER Hub

Date: 20-02-2026

Self Project Pushed in Docker Hub

The personal project was containerized and uploaded to Docker Hub, so that it can be shared, deployed, and run on any machine with Docker.

First, a Docker image was created from the project directory using a Dockerfile with this command:

```
build -t fitnessapp .
```

After the successful building of the image, it was tagged with the Docker hub username and version number, thus preparing the image for the upload:

```
docker tag fitnessapp afeefahmed/fitnessapp:v1
```

Next, the Docker Hub authentication was done by running:

```
docker login
```

Once logged in, the image was pushed to the Docker Hub repository using:

```
docker push afeefahmed/fitnessapp:v1
```

During the push operation, Docker uploaded the image layers to the Docker Hub repository. The digest message indicated that the upload was successfully completed.

This finalizes the assurance that the containerized project can be downloaded and run on any computer system, thus ensuring portability and deployment ease.




Figure 14: Self Project Pushed in Docker Hub

Pulling Self Image and Running Self Image

The project image that was previously uploaded to Docker Hub was pulled again to verify that it can be downloaded and run on any system. This is to ensure that the project image can be downloaded and run on any machine.

The image was downloaded via the below command:

```
Docker pull afeefahmed/fitnessapp:v1
```

After downloading, there was a message in Docker that the image was successfully received or was already available in the latest version.

Finally, the container was launched to run the application:

```
docker run afeefahmed/:v1
```

However, since the project is a web-based application, it was run with port mapping so that it could be accessed via a web browser:

```
docker run -p 8600:8600 afeefahmed/fitnessapp:v1
```

Once the container had started running, the application was available at: Once the container was running, the application was available under the following address:

<http://localhost:8600>

This proved that the project image was indeed pulled from Docker Hub and ran successfully, proving container portability and Docker's ability to perform rapid deployment across systems.

PERSONAL

Ctrl+K

2

A

Ask Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Images

Give feedback

Local

My Hub

59.23 GB / 30.43 GB in use

7 images

Last refresh: 1 day ago

	Name	Tag	Image ID	Created	Size	Actions
Terminal Windows PowerShell Copyright (C) Microsoft Corporation. All rights reserved. Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows PS C:\Users\afeef> docker pull afeefahmed/fitnessapp:v1 v1: Pulling from afeefahmed/fitnessapp Digest: sha256:201979e3f7cf50f952191ce0feeb7bc5110d43d91c4e3e045fd2aa44459cda2d Status: Image is up to date for afeefahmed/fitnessapp:v1 docker.io/afeefahmed/fitnessapp:v1 PS C:\Users\afeef> Windo... Windo... Windo... Windo... Windo... Windo... Windo... Engine running RAM 1.50 GB CPU 0.50% Disk: 89.19 GB used (limit 1006.85 GB) Terminal v4.62.0						

Figure 15: Pulling Self Image

docker.desktop

PERSONAL

Search

Ctrl+K

2

A

Ask Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Images

Give feedback

Local

My Hub

59.23 GB / 30.43 GB in use 7 Images

fi

Name

Tag

fitnessapp

latest

afeefahmed/fitnessapp

v1

Terminal

PS C:\Users\afeef> docker run --rm -p 8600:8501 afeefahmed/fitnessapp:v1

Collecting usage statistics. To deactivate, set browser.gatherUsageStats to false.

You can now view your Streamlit app in your browser.

Local URL: http://localhost:8501

Network URL: http://172.17.0.4:8501

External URL: http://125.214.18:8501

/app/fitness_tracker.py:173: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

sns.lineplot(

/app/fitness_tracker.py:173: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

Engine running

RAM 1.72 GB CPU 0.06% Disk: 89.20 GB used (limit 1006.85 GB)

Personal Fitness Tracker

localhost:8600

HackRx 6.0 - 2025 [...]

Refugee Allocation...

Deploy

Personal Fitness Tracker

Current User Input Summary

User Metrics

	Metric	Value
0	Age	30
1	Gender	Male
2	Weight (kg)	70.0
3	Height (m)	1.8
4	Max_BPM	150
5	Avg_BPM	120
6	Resting_BPM	60

Figure 16: Running self-image

Pulling and Running Friend Image

The docker image shared by a friend was pulled from the docker hub to verify that containers published by other users can be successfully downloaded and executed. Pulling the image ensures that a similar version of the application can be created on other machines without any need for manual installation or setup.

The image was downloaded using the following command:

```
docker pull afeefahmed/fitnessapp:v1
```

After the process was complete, a confirmation message was displayed by Docker, confirming that the image had been successfully retrieved.

Next, the application was launched by executing the container:

```
docker run afeefahmed/fitnessapp:v1
```

If the application involves a web-based interface, it can be started by port mapping to enable browser access:

```
docker run -p 8600:8600 afeefahmed/fitnessapp:v1
```

However, the application worked fine after the container was put into operation. Once the container was started, the application was running successfully, which proved that containers created by other users can be pulled and executed, showing container portability and container sharing through Docker Hub.

docker.desktop

PERSONAL

Q

Search

Ctrl+K

?

2

A

—

×

Ask Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Images

[Give feedback](#)

Local

My Hub

42.12 GB / 30.43 GB in use

7 images

Last refresh: 13 hours ago

Q

Search

		Name	Tag ↑	Image ID	Created	Size	Actions
☐	☐	python	3.10	2d8de314e8ce	20 days ago	1.6 GB	
☐	☐	websiteapp	latest	cf8fb10c9b34	9 hours ago	315.86 MB	
☐	☐	fitnessapp	latest	201979e3f7cf	1 hour ago	2.82 GB	
☐	☐	afeefahmed/fitnessapp	v1	201979e3f7cf	1 hour ago	2.82 GB	
☐	☐	sid110/hume-rainfall-app	v1	a255a3ea7b5f	10 hours ago	2.58 GB	

Showing 8 items

Terminal

+ ↓ ×

Q

Windo...

×

Windo...

×

PS C:\Users\afeef\fitness-tracker> docker pull sid110/hume-rainfall-app:v1

Error response from daemon: failed to resolve reference "docker.io/sid110/hume-rainfall-app:v1": failed to do request: Head "https://registry-1.05da006837fe: Pull complete

d85099f0969e: Pull complete

8cbc47ff628d: Pull complete

5f6e5e6733ec: Pull complete

3a3db849ea4a: Pull complete

786b7f3b9fc7: Pull complete

64faa99400e1: Pull complete

Digest: sha256:a255a3ea7b5f9ad2bd61d8a299990f2586f1911caddf307922798200f22c49a3

Status: Downloaded newer image for sid110/hume-rainfall-app:v1

docker.io/sid110/hume-rainfall-app:v1

Engine running

|| ⋮

RAM 1.73 GB CPU 0.13% Disk: 89.15 GB used (limit 1006.85 GB)

Terminal

✓ v4.62.0

Figure 17: Pulling Friend Image




Figure 18: Running Friend Image

Docker Images After All Pulls

After downloading all the necessary containers, the list of Docker images was checked to ensure that every image was successfully pulled and stored locally. Finally, the list of docker images was checked to make sure that all images have been downloaded and stored locally, which includes tensorflow, pytorch, triton server images, self project image, and friend's image.

The below command was run:

```
docker images
```

This command prints all the locally saved Docker images alongside some major details, including:

- Repository
- Tag (version)
- Image ID
- Image creation date
- Size of the image

Output images included:

- NVIDIA TensorFlow container
- NVIDIA PyTorch container
- Triton Inference Server container
- A self project image retrieved from Docker Hub (ibid).
- A friend's shared Docker image

Finally, this verification step assured that:

- All of the required images were downloaded successfully
- The images are downloaded and saved locally (Docker environment)
- The containers are in place and can be run when necessary

Hence, this confirms successfully completing the image pulling process and correctly setting up the Docker environment.




Figure 19: Docker Images After All Pulls

Task 4 : Implementing and executing Kubernetes

Date: 27-02-2026

CHECKING NODES & PODS AFTER INSTALLING MINIKUBE

After installing Minikube, a Kubernetes cluster was started to build a local single-node cluster for running containers.

The cluster was started through the command:

```
minikube start
```

This command executes the following functions:

- Creates a local Kubernetes cluster
- Sets up a virtual node using the Docker driver
- Configures kubectl communication with the cluster (ibid)

Once the startup procedure was complete without any errors, the cluster status was checked.

For checking whether the node was running correctly, the following command was fulfilled:

```
kubectl get nodes
```

The output showed the minikube node as Ready, indicating that the cluster was correctly initialized.

Next, all pods in the running status across the namespace are checked using:

```
kubectl get pods -A
```

This command outputted relevant system pods necessary for Kubernetes operations, including:

- kube-apiserver
- kube-controller-manager
- kube-scheduler
- Other core Kubernetes services

These verification steps proved that:

- Minikube was successfully started
- The Kubernetes node was in a Ready
- Verification that all system's components were in order

This made sure that the local kubernetes cluster was fully functional and ready to be used for deploying containerized applications.

PERSONAL

Search

Ctrl+K

2

A

Ask Gordon BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit BETA

Docker Hub

Docker Scout

Extensions

Images [Give feedback](#)

Local

My Hub

59.23 GB / 30.43 GB in use

7 images

Search

Name

Tag ↑

Image ID ↑

Created

Size

Actions

Terminal

Windows PowerShell

Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! <https://aka.ms/PSWindows>

PS C:\Users\afeef> kubectl get nodes

NAME	STATUS	ROLES	AGE	VERSION
minikube	Ready	control-plane	22h	v1.35.0

PS C:\Users\afeef> kubectl get pods -A

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	fitness-deployment-84d64b4b84-l668l	1/1	Running	0	12m
default	fitness-secret-deployment-5cc7dc7fd4-6lrg5	1/1	Running	0	12m
kube-system	coredns-7d764666f9-slxzbz	1/1	Running	0	12m
kube-system	coredns-7d764666f9-z4gpf	1/1	Running	0	12m
kube-system	etcd-minikube	1/1	Running	0	22h
kube-system	kube-apiserver-minikube	1/1	Running	1 (8h ago)	22h
kube-system	kube-controller-manager-minikube	1/1	Running	0	22h
kube-system	kube-proxy-mjg9n	1/1	Running	0	22h
kube-system	kube-scheduler-minikube	1/1	Running	0	22h

PS C:\Users\afeef>

Engine running

RAM 1.62 GB CPU 0.00% Disk: 89.19 GB used (limit 1006.85 GB)

Terminal ✓ v4.62.0

Figure 20: Checking Nodes & Pods After Installing

YAML DEFINING FOR SELF PULL

A Kubernetes deployment YAML file was created to deploy the self project image from Docker Hub into the Minikube cluster. This configuration file states the desired state of the application, and it includes details, such as the container image, number of replicas, and ports.

When this YAML file gets applied, Kubernetes, by default:

- Pulls the required container image from Docker Hub
- Sets up the pods to the number specified in the configuration
- Maintains the number of replicas as specified
- Guarantees that the application runs as per the configuration

The deployment configuration was saved in a file called:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: fitness-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: fitnessapp
  template:
    metadata:
      labels:
        app: fitnessapp
    spec:
      containers:
        - name: fitness-container
          image: afeefahmed/fitnessapp:v1
          ports:
            - containerPort: 8501
```

This configuration is meant to instruct Kubernetes to create a single replica of the container from self project image from Docker Hub. Upon the deployment being applied, Kubernetes downloads and executes it its cluster (on devices contained in the cluster) in case the image has not been previously downloaded. Thus, the application is run according to this configuration inside the Minikube (Kube) environment.




Figure 21: Yaml Defining For SelfPull

SELF PULL EXECUTION

After preparation of the deployment YAML file, it was deployed on the Minikube cluster to deploy the self project image. This step instructs the Kubernetes to create the pod by pulling the image from Docker Hub.

The deployment was done using the following command:

```
kubectl apply -f fitness-deployment.yaml
```

Below are the actions performed by Kubernetes upon the run of the above command:

- Created a deployment resource
- Pulled the container image from Docker Hub (if not already available)
- Started the container in the pod

To confirm that the pod was successfully created and running, the following command was applied:

```
kubectl get pods
```

The output showed one pod related to the deployment in a Running state.

This verified that the self project image was pulled successfully and then run using the Minikube Kubernetes cluster.

The screenshot displays the Docker Desktop application interface. The top navigation bar includes the Docker logo, the text "docker.desktop PERSONAL", a search bar, and a "Ctrl+K" button. The left sidebar contains a list of navigation items: Ask Gordon (BETA), Containers, Images (selected), Volumes, Kubernetes, Builds, Models, MCP Toolkit (BETA), Docker Hub, Docker Scout, and Extensions.

The main content area is divided into two sections. The top section, titled "Images", shows a progress bar for "59.23 GB / 30.43 GB in use" and "7 images". Below this is a search bar and a table with columns "Name", "Tag", and "Image ID". The bottom section, titled "Terminal", shows a Windows PowerShell session. The user has executed the following commands:

```
PS C:\Users\afeef> notepad fitness-deployment.yaml
PS C:\Users\afeef> kubectl get pods
```

The output of the `kubectl get pods` command is as follows:

NAME	READY	STATUS	RESTARTS	AGE
fitness-deployment-84d64b4b84-l668l	1/1	Running	0	25m
fitness-secret-deployment-5cc7dc7fd4-6lrg5	1/1	Running	0	25m

The bottom status bar indicates "Engine running", "RAM 1.66 GB", "CPU 0.06%", "Disk: 89.19 GB used (limit 1006.85 GB)", and "v4.62.0".

Overlaid on the right side of the terminal is a text editor window titled "fitness-deployment.yaml". It contains the following YAML content:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: fitness-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: fitnessapp
  template:
    metadata:
      labels:
        app: fitnessapp
    spec:
      containers:
        - name: fitness-container
          image: afeefahmed/fitnessapp:v1
          ports:
            - containerPort: 8501
```

Figure 22: Self Pull Execution

UPSCALE

After deployment of the application in Minikube was successful, the deployment was scaled up to increase the number of running pods. Scaling enables Kubernetes to run multiple instances of the same application, which is useful for:

- Cater for more users
- Increase reliability through redundancy
- Enhance the availability of the application

The deployment was scaled by the following command:

```
kubectrl scale deployment fitness-deployment --replicas=3
```

This command instructs Kubernetes to run three replicas of the container instead of one. To scale out the container successfully, the following command was run to ascertain the same:

```
kubectrl get pods
```

The output showed three pods in the Running state, all linked to the deployment.

This verified that the application was scaled successfully and multiple instances were running within Minikube cluster.




Figure 23: UPSCALE

DOWNSCALE

After increasing the number of replicas, the deployment was scaled down to demonstrate a Kubernetes downscaling. Downscaling also achieves a reduction in the number of active pods, thus releasing unnecessary system resources.

The deployment was downscaled by running the below command:

```
kubectrl scale deployment fitness-deployment --replicas=1
```

This command tells Kubernetes to scale down the number of replicas to one.

To ensure that the downscaling operation was carried out successfully, the following command was run:

```
kubectrl get pods
```

There was only one Running pod in the output, indicating that the deployment had been scaled down to one replica as intended.

This demonstrated Kubernetes' ability to dynamically create the number of instances of an application required at any given time.

docker.desktop

PERSONAL

Q Search

Ctrl+K

2

A

Ask Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Images

[Give feedback](#)

Local

My Hub

59.23 GB / 30.43 GB in use

7 images

Last refresh: 1 day ago

Q Search

	Name	Tag	Image ID	Created	Size	Actions
--	------	-----	----------	---------	------	---------

Terminal

Windows PowerShell

Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\afeef> kubectl scale deployment fitness-deployment --replicas=1

deployment.apps/fitness-deployment scaled

PS C:\Users\afeef> kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
fitness-deployment-84d64b4b84-l668l	1/1	Running	0	34m
fitness-secret-deployment-5cc7dc7fd4-6lrg5	1/1	Running	0	34m

PS C:\Users\afeef>

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Engine running

RAM 1.65 GB CPU 0.00% Disk: 89.19 GB used (limit 1006.85 GB)

> Terminal

v4.62.0

Figure 24: Downscale

SECRET YAML DEFINING

A Kubernetes Secret YAML file was created to hold the configuration data securely within the Minikube cluster. Secrets are useful to store passwords, API keys, or environmental variables without hardcoding them into the application.

The secret configuration was saved into a file called:

```
apiVersion: v1
kind: Secret
metadata:
  name: fitness-secret
type: Opaque
data:
  username: YWRtaW4=
  password: MTIzNA==
```

Here is the configuration:

- A Secret named app-secret is created
- The type Opaque defines a generic secret
- The key-value pair has APP_MODE set to production
- The stringData field accepts non-encoded data, which Kubernetes re-encodes using Base64 internally

This secret can later be referenced in a deployment file, usually as an environment variable in a container. This approach secures the sensitive configuration with a clear distinction from application code.

The screenshot displays the Docker Desktop application interface. The top navigation bar includes the Docker logo, a search bar, and various utility icons. The left sidebar contains a list of Docker features: Ask Gordon (BETA), Containers, Images (selected), Volumes, Kubernetes, Builds, Models, MCP Toolkit (BETA), Docker Hub, Docker Scout, and Extensions.

The main content area is divided into two sections. The top section, titled "Images", shows a progress bar indicating "59.23 GB / 30.43 GB in use" and "7 Images". Below this is a search bar and a table with columns for Name, Tag, and Image ID. The bottom section, titled "Terminal", shows a Windows PowerShell session. The user has executed the following commands:

```
PS C:\Users\afeef> kubectl scale deployment fitness-deployment --replicas=1
deployment.apps/fitness-deployment scaled
PS C:\Users\afeef> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
fitness-deployment-84d64b4b84-l668l	1/1	Running	0	34m
fitness-secret-deployment-5cc7dc7fd4-6lrg5	1/1	Running	0	34m

```
PS C:\Users\afeef> notepad secret.yaml
PS C:\Users\afeef>
```

The "secret.yaml" file is open in a text editor, showing the following content:

```
apiVersion: v1
kind: Secret
metadata:
  name: fitness-secret
type: Opaque
data:
  username: YWRtaW4=
  password: MTIzNA==
```

The bottom status bar indicates "Engine running" and provides system resource usage: "RAM 1.66 GB CPU 0.00% Disk: 89.19 GB used (limit 1006.85 GB)". The version "v4.62.0" is also displayed.

Figure 25: Secret Yaml Defining

NEW DEPLOYMENT DEFINING KIND SECRET

To show how a Secret can be passed into an app via environment variables, a new Kubernetes deployment configuration was done. This deployment injects Secret value into a running container instead of hardcoding it into an application.

The configuration was saved in a file called:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: fitness-secret-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: fitnesssecret
  template:
    metadata:
      labels:
        app: fitnesssecret
    spec:
      containers:
        - name: fitness-container
          image: afeefahmed/fitnessapp:v1
          ports:
            - containerPort: 8501
          env:
            - name: USERNAME
              valueFrom:
                secretKeyRef:
                  name: fitness-secret
                  key: username
            - name: PASSWORD
              valueFrom:
                secretKeyRef:
                  name: fitness-secret
                  key: password
```

Explanation of the Configuration

- A deployment named fitness-secret-deployment is created
- It runs a single replica container based on the self project image from Docker Hub.
- The selector and labels are in place to ensure the proper identification and handling of the associated pods by Kubernetes.
- The vital part is the env section, where:
 - There is a set definition for the environment variable APP_MODE.
 - It is not specified directly in the file..
 - However, the value is not put directly in the file but retrieved securely from the Kubernetes Secret called app-secret.
 - The secretKeyRef defines the secret name and the particular key necessary for access.




Figure 26: New Deployment Defining Kind Secret

GETTING SECRET FOR CHECK

Once the Kubernetes Secret had been generated, its presence and contents were validated by using various commands in order to prove that the Secret had been correctly stored within the Kubernetes cluster. Several commands can be used to inspect Kubernetes Secrets.

In order to list all the secrets stored within the cluster, the following command was run:

```
kubectrl get secrets
```

As can be seen in the output above, there is indeed a secret called app-secret, which has been created.

In order to get additional information regarding the metadata associated with this secret, the following command was run:

```
kubectrl describe secret app-secret
```

Additional metadata obtained from this secret include:

- Secret name
- The type of secret
- The number of keys stored in it

Finally, in order to validate that the entire content of this secret had been correctly stored, the following command was run:

```
kubectrl get secret app-secret -o yaml
```

The stored values are shown in Base64 encoded format. In order to actually get the value of the APP_MODE field, the following command was used:

```
kubectrl get secret app-secret -o jsonpath="{.data.APP_MODE}" | base64 -- decode
```

Conclusion

The verification steps shown in this output demonstrate that:

- The Kubernetes secret has been successfully created
- It has been securely stored within the Kubernetes cluster
- The values stored within it can be accessed when necessary




Figure 27: Getting Secret For Check

SECRET DEPLOYMENT

Once the Kubernetes secret had been created, its creation and storage were validated in order to demonstrate proper creation. There are many commands available within Kubernetes that could verify the secret in question.

Firstly, the following command was used in order to list all the secrets currently available within the Kubernetes cluster:

```
kubectrl get secrets
```

It shows that there is a secret called app-secret available within the cluster.

The following command was used in order to display additional metadata:

```
kubectrl describe secret app-secret
```

Additional metadata revealed by this secret include:

- Secret name
- Secret type (Opaque)
- Keys stored within the secret

In order to verify the data stored within the secret, the following command was used:

```
kubectrl get secret app-secret -o yaml
```

Note that the stored values are in Base64 format, because Kubernetes encrypts the secret data automatically.

```
kubectrl get secret app-secret -o jsonpath="{.data.APP_MODE}" | base64 -- decode:
```

```
kubectrl get secret app-secret -o jsonpath="{.data.APP_MODE}" | base64 -- decode
```

Conclusion

The verification steps demonstrated above confirm that:

- The Kubernetes secret has been successfully created
- It has been securely stored within the Kubernetes cluster
- The values stored within it can be accessed whenever required

docker.desktop

PERSONAL

Q

Search

Ctrl+K

?

2

A

—

×

✦ Ask Gordon BETA

📦 Containers

🔄 Images

💾 Volumes

⚙️ Kubernetes

🔧 Builds

🔗 Models

🔗 MCP Toolkit BETA

🌐 Docker Hub

🔍 Docker Scout

🔗 Extensions

Images

[Give feedback](#)

Local

My Hub

59.23 GB / 30.43 GB in use

7 images

Last refresh: 1 day ago

🔄

Q

Search

≡

☰

	Name	Tag	Image ID	Created	Size	Actions
--	------	-----	----------	---------	------	---------

Terminal

secret/app-secret created

PS C:\Users\afeef> kubectl get secrets

NAME TYPE DATA AGE

app-secret Opaque 1 5s

fitness-secret Opaque 2 23h

PS C:\Users\afeef> kubectl get secret app-secret -o yaml

apiVersion: v1

data:

APP_MODE: cHJvZHVjdGlvbg==

kind: Secret

metadata:

creationTimestamp: "2026-02-26T08:07:22Z"

name: app-secret

namespace: default

resourceVersion: "9356"

uid: e7554792-8040-4195-b23a-a892507230cf

type: Opaque

PS C:\Users\afeef> kubectl get secret app-secret -o jsonpath="{.data.APP_MODE}" | %[[System.Text.Encoding]::UTF8.GetString([System.Convert]::FromBase64String(\$_))]

production

PS C:\Users\afeef> kubectl get pods

NAME READY STATUS RESTARTS AGE

fitness-deployment-84d64b4b84-l668l 1/1 Running 0 53m

fitness-secret-deployment-5cc7dc7fd4-6lrg5 1/1 Running 0 53m

PS C:\Users\afeef>

🔍

📄

🗑️

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Windo...

Engine running

RAM 1.65 GB CPU 0.12% Disk: 89.19 GB used (limit 1006.85 GB)

Terminal

v4.62.0

Figure 28: Secret Deployment

UPSCALE

In order to further demonstrate the capabilities of the application and the platform, the deployment was upscaled to increase the number of active pods. Upscaling makes the application more reliable and proves that the Kubernetes platform can handle increasing workloads.

The deployment is upscaled using the following command:

```
kubectl scale deployment fitness-secret-deployment --replicas=3
```

This instruction tells Kubernetes to make three copies of the pods related to this deployment.

To verify whether the deployment is upscaled successfully, the following command is executed:

```
kubectl get pods
```

As a result, three pods with the status of Running are observed with the name of fitness-secret-deployment.




Figure 29: UPSCALE

DOWNSCALE

Having demonstrated the upscaled deployment, the deployment was downscaled, which means that fewer pods will be running.

Downscaling helps optimize the usage of resources and prove that Kubernetes can effectively adjust its capacity for an application.

The deployment is downscaled using the following command:

```
kubectrl scale deployment fitness-secret-deployment --replicas=1
```

```
kubectrl scale deployment fitness-secret-deployment --replicas=1
```

This command instructs Kubernetes to scale down to one copy of the deployment.

To verify whether the operation was successful, the following command is executed:

```
kubectrl get pods
```

As a result, one pod with the status of Running is observed.

This demonstrates Kubernetes' ability to efficiently manage scaling operations while maintaining application stability.




Figure 30: Downscaling

ALL FINAL DEPLOYMENT

Upon successfully completing the process of deployment, performing scaling operations, and adding secrets to the application, it is important to check all of the deployments within the cluster. This will help ensure that the applications are successfully deployed in the cluster.

The following command will show all of the deployments within the cluster:

```
kubectl get deployments
```

The output will show the following information regarding the deployment:

- Name of the deployment
- Desired number of pods
- Available pods
- status

There are deployments such as:

- **Fitness deployment**
- **Fitness secret deployment**
- **Hume rainfall deployment**

All of these deployments have the necessary number of pods running in the READY state.

To receive more information about any deployment, the following command may be executed:

```
kubectl describe deployment fitness-secret-deployment
```

The output will show:

- Information about the pod template
- List of replica sets
- Scaling events and updates.




Figure 31: All Final Deployment

ALL PODS AT THE LAST

Once all deployment, scaling, and secret operations have been completed, the status of all pods in the Minikube cluster was then assessed to determine whether the applications have run smoothly and have been managed effectively.

Listing all running pods is made possible using the following command:

```
kubectl get pods
```

The following are some features displayed by the output of this command:

- Name of the Pod
- Readiness of the Pod
- Status of the Pod
- Number of restarts of the Pod
- Age of the Pod

Examples of pods include those associated with:

- **fitness-deployment**
- **fitness-secret-deployment**
- **hume-rainfall-deployment**

All the pods in this assessment had the status “Running,” with the correct number of running containers (e.g., 1/1).

docker.desktop

PERSONAL

Q Search

Ctrl+K

?

2

A

Ask Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Images

Give feedback

Local

My Hub

61.22 GB / 30.43 GB in use 7 images

Last refresh: 2 days ago

Q Search

Name

Tao

Image ID

Created

Size

Actions

Terminal

Windows PowerShell

Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\afeef> kubectl get deployments

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
fitness-deployment	1/1	1	1	23h
fitness-secret-deployment	1/1	1	1	23h
hume-rainfall-deployment	1/1	1	1	4m44s

PS C:\Users\afeef>

RAM 1.55 GB CPU 1.00% Disk: 89.19 GB used (limit 1006.85 GB)

>_ Terminal

v4.62.0

Figure32: AllPods

Task 5 : Working with SLURM - Scripting & Kubeflow

Date: 27-03-2026

PART 1 — SLURM

Procedure

The purpose of this study was to analyze how job scheduling is performed using the SLURM workload manager in a container environment. To perform the experiment, the SLURM cluster was created based on the Docker platform that comprised a controller node as well as compute nodes. In order to test SLURM functionality, a batch job consisting of matrix addition implemented with Python was run.

Step-by-Step Explanation

The SLURM cluster setup was done using Docker-based cluster setup.

```
git clone https://github.com/giovtorres/slurm-docker-cluster.git
cd slurm-docker-cluster
docker compose up -d
```

These clusters running were checked by doing the following commands:

```
docker ps
docker exec -it slurmctld bash
sinfo
```

A simple Python script to add two matrices:

```
A = [[1, 2], [3, 4]]
B = [[5, 6], [7, 8]]

result = [[A[i][j] + B[i][j] for j in range(2)] for i in range(2)]
print(result)
```

PERSONAL

Ctrl+K

FPS N/A | GPU 0% 52°C | CPU 4% | LAT N/A

A

Gordon BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit BETA

Docker Hub

Docker Scout

Extensions

Containers [Give feedback](#)

Container CPU usage

0.43% / 1600% (16 CPUs available)

Container memory usage

175.8MB / 7.39GB

☐ Only show running containers

		Name	Container ID	Image	Port(s)
☐	☐	charming_easley	9383cb167579	sid110/hume-rainfall-app	
☐	☐	amazing_nightingale	452b3aeb7293	nvidia/pytorch:26.01-py3	

Terminal

```

✓ Volume slurm-docker-cluster_var_log_slurm Created 10.8s
✓ Volume slurm-docker-cluster_home_ood Created 2.5s
✓ Volume slurm-docker-cluster_etc_munge Created 0.6s
✓ Volume slurm-docker-cluster_opt_modulefiles Created 0.1s
90f7cca4127f slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 29 seconds (healthy) 6818/tcp
slurm-docker-cluster-cpu-worker-1
4ce3b1311ba0 slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 29 seconds (healthy) 0.0.0.0:6820->6820/tcp, [::]:6820->6820/tcp
slurmrestd
602e901d9e50 slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 34 seconds (healthy) 0.0.0.0:3022->22/tcp, [::]:3022->22/tcp
slurmctld
7de6de14ede0 slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 40 seconds (healthy) 6819/tcp
slurmdbd
6ee4df25bf17 mariadb:12 "docker-entrypoint.s..." 53 seconds ago Up 51 seconds (healthy) 3306/tcp
mysql
PS C:\Users\afeef\slurm-docker-cluster> docker exec -it slurmctld bash
[root@slurmctld data]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE MODELIST
cpu*      up    infinite    1  idle c1
gpu       up    infinite    0   n/a
[root@slurmctld data]#

```

Upgrade plan

Engine running

RAM 7.25 GB CPU 6.41% Disk: 123.13 GB used (limit 1006.85 GB)

A

PERSONAL

afeefahmed

What's new

Account Settings

Upgrade

Sign out

Show charts

⋮

🗑

⋮

🗑

+

⌵

×

Windo...

×

Terminal

✓ v4.71.0

Fig 33: Docker Containers Running (SLURM Cluster)

PERSONAL

Ctrl+K

?

1

✉

⚙

⌵

A

Gordon BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit BETA

Docker Hub

Docker Scout

Extensions

Containers [Give feedback](#)

Container CPU usage ?

0.43% / 1600% (16 CPUs available)

Container memory usage ?

175.8MB / 7.39GB

☰

Only show running containers

☐		Name	Container ID	Image	Port(s)
☐	☐	charming_easley	9383cb167579	sid110/hume-rainfall-app:	
☐	☐	amazing_nightingale	452b3aeb7293	nvidia/pytorch:26.01-py3	

Terminal

```

✓ Volume slurm-docker-cluster_var_log_slurm Created 10.8s
✓ Volume slurm-docker-cluster_home_ood Created 2.5s
✓ Volume slurm-docker-cluster_etc_munge Created 0.6s
✓ Volume slurm-docker-cluster_opt_modulefiles Created 0.1s
90f7cca4127f slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 29 seconds (healthy) 6818/tcp
slurm-docker-cluster-cpu-worker-1
4ce3b1311ba0 slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 29 seconds (healthy) 0.0.0.0:6820->6820/tcp, [::]:6
820->6820/tcp slurmrestd
602e901d9e50 slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 34 seconds (healthy) 0.0.0.0:3022->22/tcp, [::]:302
2->22/tcp slurmctld
7de6de14ede0 slurm-docker-cluster:25.11.4 "/usr/local/bin/dock..." 53 seconds ago Up 40 seconds (healthy) 6819/tcp
slurmdbd
6ee4df25bf17 mariadb:12 "docker-entrypoint.s..." 53 seconds ago Up 51 seconds (healthy) 3306/tcp
mysql
PS C:\Users\afeef\slurm-docker-cluster> docker exec -it slurmctld bash
[root@slurmctld data]# sinfo
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST
cpu* up infinite 1 idle c1
gpu up infinite 0 n/a
[root@slurmctld data]#

```

Upgrade plan

Engine running

RAM 7.25 GB CPU 6.41% Disk: 123.13 GB used (limit 1006.85 GB)

A

PERSONAL

afeefahmed

What's new

Account Settings

Upgrade

Sign out

Show charts

⋮

🗑

⋮

🗑

+

✓

✕

Windo...

✕

Terminal

✓ v4.71.0

Fig 34: SLURM Cluster Status (sinfo)

A batch script with SLURM commands was developed as follows:

```
#!/bin/bash
#SBATCH --job-name=matrix_add
#SBATCH --output=result.txt
python3 matrix.py
```

The script was successfully run:

```
sbatch matrix.slurm
cat result.txt
```

As seen from the result, the experiment succeeded as SLURM was able to schedule and execute the command in question.

PART 2 — Kubeflow

Procedure

The goal of the experiment was workflow automation with Kubeflow Pipelines. Thus, a matrix addition component had to be created and compiled into a YAML file describing the workflow process.

Step-by-Step Explanation

Firstly, we installed Kubeflow Pipelines SDK to allow us to define our workflows:

```
pip install kfp
```




Fig 38:Kubeflow Pipeline Script

A script in Python was developed to implement a reusable module to perform matrix addition. This involved building logic for matrix addition:

```
from kfp import dsl
from kfp.compiler import Compiler

@dsl.component
def matrix_addition():
    print("Matrix Addition using Kubeflow")

    A = [[1, 2], [3, 4]]
    B = [[5, 6], [7, 8]]

    result = [[A[i][j] + B[i][j] for j in range(2)] for i in range(2)]

    print("Result:")
    for row in result:
        print(row)

@dsl.pipeline(name="matrix-addition-pipeline")
def pipeline():
    matrix_addition()
```

Execution:

```
python pipeline.py
```

The resulting pipeline.yaml indicated that the pipeline was successfully created..

Task 6 : Using Jupyter Notebook and PyTorch for Tensor Operations and Data-Centric Application

Date: 08-04-2026

Experiment Procedure

This experiment was designed to understand tensor operations and create a data-centric application using Jupyter Notebook and PyTorch. The following steps were executed to complete the task.

Step-by-Step Explanation

Installation of required libraries and opening of notebook.

Performance of tensor operations using Python script.

Creation of data-centric application using linear regression model.

Experiment Steps Explanation

In order to perform the experiment, all the required libraries were installed and Jupyter Notebook was opened:

```
pip install notebook torch torchvision
jupyter notebook
```

After opening the notebook, two tensors were created and several operations like addition, subtraction, element-wise multiplication, and matrix multiplication were executed on them.

```
import torch
```

```
a = torch.tensor([[1, 2], [3, 4]])
b = torch.tensor([[5, 6], [7, 8]])
```

```
print("Addition:\n", a + b)
print("Subtraction:\n", a - b)
print("Element-wise Multiplication:\n", a * b)
print("Matrix Multiplication:\n", torch.matmul(a, b))
```

Further, the data-centric application was prepared by creating a dataset containing linearly related variables and defining a linear model using PyTorch.

```
import torch.nn as nn
import torch.optim as optim
```

```
x = torch.tensor([[1.0], [2.0], [3.0], [4.0]])
y = torch.tensor([[3.0], [5.0], [7.0], [9.0]])
```

```
model = nn.Linear(1, 1)
```





Fig 40:Jupyter Installation Process



The screenshot shows a JupyterLab interface in a web browser. The browser's address bar displays `localhost:8888/notebooks/pytorch_task.ipynb`. The JupyterLab header includes the logo, the notebook name `pytorch_task`, and a status message `Last Checkpoint: 2 minutes ago`. Below the header is a menu bar with `File`, `Edit`, `View`, `Run`, `Kernel`, `Settings`, and `Help`. A toolbar contains icons for file operations and execution. The main area shows a code cell with the following Python code:

```
# Addition
print("\nAddition:\n", a + b)
# Subtraction
print ("\nSubtraction:\n", a - b)
# Multiplication
print ("\nElement-wise Multiplication:\n", a * b)
# Matrix Multiplication
print ("\nMatrix Multiplication:\n", torch.matmul(a, b))
```

The output of the code cell is displayed below the code:

```
Basic Tensor Operations
Tensor A:
tensor([[1, 2],
        [3, 4]])
Tensor B:
tensor([[5, 6],
        [7, 8]])


Addition:
tensor([[ 6,  8],
        [10, 12]])

Subtraction:
tensor([[ -4,  -4],
        [-4,  -4]])

Element-wise Multiplication:
tensor([[ 5, 12],
        [21, 32]])

Matrix Multiplication:
tensor([[19, 22],
        [43, 50]])
```

Fig 41:Jupyter Tensor Output



The screenshot shows a JupyterLab notebook titled 'pytorch_task' running on a local host. The notebook contains a Python script for training a linear regression model using PyTorch. The script defines sample data, a linear model, a loss function, and an optimizer. It then enters a training loop for 100 epochs, calculating the loss and updating the model parameters. After training, it prints the trained model parameters and makes a prediction for x=5.

```
# Sample data (y = 2x + 1)
x = torch.tensor([[1.0], [2.0], [3.0], [4.0]])
y = torch.tensor([[3.0], [5.0], [7.0], [9.0]])
# Model
model = nn.Linear(1, 1)
# Loss and optimizer
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
# Training Loop
for epoch in range(100):
    pred = model(x)
    loss = criterion(pred, y)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
print("\nTrained Model:")
for name, param in model.named_parameters():
    print(name, param.data)
# Test prediction
test = torch.tensor([[5.0]])
print("\nPrediction for x=5:", model(test).item())
```

Simple Data-Centric Application (Linear Regression)

Trained Model:
weight tensor([[1.8377]])

Prediction for x=5: 10.665677070617676
bias tensor([1.4772])

Prediction for x=5: 10.665677070617676

Fig 42:Linear Regression Output

After preparing the dataset and defining the model, the next step involved specifying a loss function and an optimizer in order to train the linear regression model iteratively.

```
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

```
for epoch in range(100):
    pred = model(x)
    loss = criterion(pred, y)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

After training the model using iterative epochs, the final step was to test the trained model using test inputs.

```
test = torch.tensor([[5.0]])
print("Prediction for x=5:", model(test).item())
```

Finally, we checked whether the newly trained linear model gave correct predictions for new inputs.

Task 7 : Testing the GPGPU Computing

Date: 21-04-2026

Procedure

In this experiment, analysis and comparison of CPU-based computation vs GPU-based computation was carried out using PyTorch. The objective was to show the improvements offered by the use of GPUs for parallel computing on large scale tasks.

Step-by-Step Explanation

In order to confirm the availability of a GPU, the following code snippet was executed.

```
import torch

print("GPU Available:", torch.cuda.is_available())
```

Firstly, a large dataset was created. Sorting operation was performed using CPU, and runtime was measured.

```
import time

data = torch.rand(10_000_000)

start = time.time()
torch.sort(data)
cpu_time = time.time() - start

print("CPU Time:", cpu_time)
```

The screenshot displays a Jupyter Notebook environment within a web browser. The browser's address bar shows a Google Drive link. The notebook interface includes a top menu bar with options like File, Edit, View, Insert, Runtime, Tools, and Help. Below the menu is a toolbar with icons for commands, code, text, and a 'Run all' button. The notebook itself contains two code cells. The first cell, labeled [1], has a green checkmark and a 4s execution time. It contains Python code to import torch and print GPU availability and device name. The output shows 'GPU Available: True' and 'Device: Tesla T4'. The second cell, labeled [2], also has a green checkmark and a 2s execution time. It contains code to create a large dataset, perform CPU and GPU sorting, and print the sorting times. The output shows 'CPU Sorting Time:' and 'GPU Sorting Time:'. The bottom status bar indicates the current time as 4:14 AM and the environment as T4 (Python 3).

```
[1] ✓ 4s
import torch
print("GPU Available:", torch.cuda.is_available())
print("Device:", torch.cuda.get_device_name(0))

... GPU Available: True
Device: Tesla T4

[2] ✓ 2s
import torch
import time

# Create large dataset
size = 10_000_000
data = torch.rand(size)

# ----- CPU SORT -----
start = time.time()
cpu_sorted = torch.sort(data)
cpu_time = time.time() - start

print("CPU Sorting Time:", cpu_time)

# ----- GPU SORT -----
device = torch.device("cuda")

data_gpu = data.to(device)

start = time.time()
gpu_sorted = torch.sort(data_gpu)
torch.cuda.synchronize() # important for accurate timing
gpu_time = time.time() - start

print("GPU Sorting Time:", gpu_time)
```

Fig 43:GPU Verification Output

The screenshot displays a JupyterLab environment with a web browser at the top showing the URL `colab.research.google.com/drive/107B18_Sns-jsLLZbSgH3OgwsvjG7z7Ry#scrollTo=P5QUbY8Ospa-`. The JupyterLab interface includes a menu bar (File, Edit, View, Insert, Runtime, Tools, Help), a toolbar (Commands, Code, Text, Run all), and a sidebar with icons for file management and settings. The main workspace shows two code cells. The first cell, labeled [1], contains the code `print("Device:", torch.cuda.get_device_name(0))` and its output is `GPU Available: True` and `Device: Tesla T4`. The second cell, labeled [2], contains a more complex script for sorting data on CPU and GPU. The script defines a large dataset of size 10,000,000, sorts it on the CPU, and then sorts it on the GPU. The output of the second cell shows the CPU sorting time as 1.619107723236084 and the GPU sorting time as 0.06518912315368652. The bottom status bar indicates the current time is 4:14 AM and the device is T4 (Python 3).

```
[1] ✓ 4s
print("Device:", torch.cuda.get_device_name(0))

... GPU Available: True
Device: Tesla T4

[2] ✓ 2s
import torch
import time

# Create large dataset
size = 10_000_000
data = torch.rand(size)

# ----- CPU SORT -----
start = time.time()
cpu_sorted = torch.sort(data)
cpu_time = time.time() - start

print("CPU Sorting Time:", cpu_time)

# ----- GPU SORT -----
device = torch.device("cuda")

data_gpu = data.to(device)

start = time.time()
gpu_sorted = torch.sort(data_gpu)
torch.cuda.synchronize() # important for accurate timing
gpu_time = time.time() - start

print("GPU Sorting Time:", gpu_time)

... CPU Sorting Time: 1.619107723236084
GPU Sorting Time: 0.06518912315368652
```

Fig 44:GPU Sorting Output

Then, the same data was moved to the GPU memory, and the sorting operation was again performed to measure runtime.

```
device = torch.device("cuda")
```

```
data_gpu = data.to(device)
```

```
start = time.time()
```

```
torch.sort(data_gpu)
```

```
torch.cuda.synchronize()
```

```
gpu_time = time.time() - start
```

```
print("GPU Time:", gpu_time)
```

As seen from the results, GPU computations were considerably more efficient compared to CPU for large data sets, which demonstrated the benefits of parallel computations.

Task 8 : Testing CUDA-X Tensor Operations, Mixed Precision, and Quantization-Aware Training using PyTorch

Date: 25-04-2026

Procedure

In the current experiment, we explored GPU programming in PyTorch through advanced CUDA tensor operations, mixed precision, and automatic differentiation.

Step-by-Step Explanation

In the following example, the device was selected as either GPU or CPU based on availability.

```
import torch

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using:", device)
```

Large tensors were created on GPU and matrix multiplication was applied to them.

```
A = torch.randn(4096, 4096, device=device)
B = torch.randn(4096, 4096, device=device)

C = A @ B
print("Matrix multiplication completed")
```

Autodiff capabilities in PyTorch were tested next.

```
x = torch.tensor(2.0, requires_grad=True)
y = x**3 + 2*x**2 + x

dy_dx = torch.autograd.grad(y, x, create_graph=True)[0]
d2y_dx2 = torch.autograd.grad(dy_dx, x)[0]

print("First derivative:", dy_dx.item())
print("Second derivative:", d2y_dx2.item())
```

Task Assistance Request afeef(task-8).ipynb - Col Home Task Assistance Request SEAL - Virtualization, Do +

colab.research.google.com/drive/107B18_Sns-jsLLZbSgH3OgwsvjG7z7Ry#scrollTo=VmhmeOptDxi

HackRx 6.0 - 2025 [...] Refugee Allocation... Research paper in IE... Quotex trading plat... pg.phaseganging.sh... Home - NetMirror FlickyStream - Watc... Vidbox - Your Ultim...

afeef(task-8).ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

RAM Disk

[3] ✓ 0s

```
import torch, time
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using:", device)
if device.type == "cuda":
    print(torch.cuda.get_device_name(0))
```

... Using: cuda
Tesla T4

[4] ✓ 0s

```
# Large tensors on GPU
n = 4096
A = torch.randn(n, n, device=device)
B = torch.randn(n, n, device=device)

# Warm-up (important for fair timing)
_ = A @ B
if device.type == "cuda":
    torch.cuda.synchronize()

# Time matrix multiplication
start = time.time()
C = A @ B
if device.type == "cuda":
    torch.cuda.synchronize()
t_gpu = time.time() - start

print("GPU matmul time:", t_gpu, "seconds")
```

GPU matmul time: 0.04804396629333496 seconds

[5] ✓ 0s

```
# First & second derivatives using autograd
x = torch.tensor(2.0, requires_grad=True)
```

Variables Terminal

4:14 AM T4 (Python 3)

Fig 45:GPU Tensor Operations




Fig 46:Autograd Derivatives Output

The screenshot displays a Jupyter Notebook titled 'afeef(task-8).ipynb' in a web browser. The notebook contains two code blocks. The first block, labeled [6], shows a standard PyTorch training loop with a for loop of 20 iterations, calculating loss and performing backpropagation. The second block, labeled [7], shows the same training loop but with AMP (Automatic Mixed Precision) enabled using `torch.cuda.amp.autocast`. The output of the notebook shows the execution time for both methods.

```
start = time.time()
for _ in range(20):
    optimizer.zero_grad()
    out = model(x)
    loss = ((out - y)**2).mean()
    loss.backward()
    optimizer.step()
if device.type == "cuda":
    torch.cuda.synchronize()
t_fp32 = time.time() - start

# AMP training
start = time.time()
for _ in range(20):
    optimizer.zero_grad()
    with torch.cuda.amp.autocast(enabled=(device.type=="cuda")):
        out = model(x)
        loss = ((out - y)**2).mean()
    scaler.scale(loss).backward()
    scaler.step(optimizer)
    scaler.update()
if device.type == "cuda":
    torch.cuda.synchronize()
t_amp = time.time() - start

print("FP32 time:", t_fp32)
print("AMP time:", t_amp)
```

Output:

```
... /tmp/ipykernel_2024/3546142149.py:8: FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated. Please use `torch.amp.GradScaler('cuda', args...)` instead.
scaler = torch.cuda.amp.GradScaler(enabled=(device.type=="cuda"))
/tmp/ipykernel_2024/3546142149.py:29: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
with torch.cuda.amp.autocast(enabled=(device.type=="cuda")):
FP32 time: 0.4018263816833496
AMP time: 0.23004746437072754
```

Fig 47:AMP Performance Output

The screenshot displays a Jupyter Notebook environment with the following components:

- Browser Tabs:** Task Assistance Request, afeef(task-8).ipynb - Colab, Home, Task Assistance Request, SEAI - Virtualization, Do...
- Address Bar:** colab.research.google.com/drive/107B18_Sns-jsLLZbSgH3OgwsvjG7z7Ry#scrollTo=VmhmeOptDxi
- File Bar:** afeef(task-8).ipynb, File, Edit, View, Insert, Runtime, Tools, Help
- Code Editor:**

```
torch.nn.Linear(10, 1)

model_q.qconfig = tq.get_default_qat_qconfig("fbgemm")
tq.prepare_qat(model_q, inplace=True)

# Fake training step
for _ in range(5):
    data = torch.randn(4, 10)
    out = model_q(data)

# Convert to quantized model
model_q.eval()
tq.convert(model_q, inplace=True)

print("Quantized model ready")
```
- Output:**

```
*** Quantized model ready
/tmp/ipykernel_2024/1969022726.py:10: DeprecationWarning: torch.ao.quantization is deprecated and will be removed in 2.10.
For migrations of users:
1. Eager mode quantization (torch.ao.quantization.quantize, torch.ao.quantization.quantize_dynamic), please migrate to use torchao eager mode quantize_ API instead
2. FX graph mode quantization (torch.ao.quantization.quantize_fx.prepare_fx, torch.ao.quantization.quantize_fx.convert_fx, please migrate to use torchao pt2e quantization API instead (prepare
3. pt2e quantization has been migrated to torchao (https://github.com/pytorch/ao/tree/main/torchao/quantization/pt2e)
see https://github.com/pytorch/ao/issues/2259 for more details
tq.prepare_qat(model_q, inplace=True)
/usr/local/lib/python3.12/dist-packages/torch/ao/quantization/observer.py:534: UserWarning: Please use quant_min and quant_max to specify the range for observers.
    super().__init__(
/tmp/ipykernel_2024/1969022726.py:19: DeprecationWarning: torch.ao.quantization is deprecated and will be removed in 2.10.
For migrations of users:
1. Eager mode quantization (torch.ao.quantization.quantize, torch.ao.quantization.quantize_dynamic), please migrate to use torchao eager mode quantize_ API instead
2. FX graph mode quantization (torch.ao.quantization.quantize_fx.prepare_fx, torch.ao.quantization.quantize_fx.convert_fx, please migrate to use torchao pt2e quantization API instead (prepare
3. pt2e quantization has been migrated to torchao (https://github.com/pytorch/ao/tree/main/torchao/quantization/pt2e)
see https://github.com/pytorch/ao/issues/2259 for more details
    tq.convert(model_q, inplace=True)
```
- Bottom Bar:** Variables, Terminal, 4:14 AM, T4 (Python 3)

Fig 48:Quantization Output

The screenshot shows a Jupyter Notebook interface with a dark theme. The browser address bar displays a Google Drive link. The notebook's top bar includes the file name 'afeef(task-8).ipynb' and a 'Share' button. The left sidebar contains navigation icons. The main area shows a code cell with a warning about deprecated quantization APIs and a code block for performance testing. The output of the code block is visible below it.

Warning: `torch.quantization` is deprecated and will be removed in 2.10. For migrations of users:

1. Eager mode quantization (`torch.quantization.quantize`, `torch.quantization.quantize_dynamic`), please migrate to use `torchao` eager mode `quantize` API instead
2. FX graph mode quantization (`torch.quantization.quantize_fx.prepare_fx`, `torch.quantization.quantize_fx.convert_fx`), please migrate to use `torchao` pt2e quantization API instead (prepare see <https://github.com/pytorch/ao/issues/2259> for more details
3. pt2e quantization has been migrated to `torchao` (<https://github.com/pytorch/ao/tree/main/torchao/quantization/pt2e>)

```
tq.convert(model_q, inplace=True)
```

```
[8] print("\n--- Performance Summary ---")
✓ 0s print("FP32 Time:", t_fp32)
    print("AMP Time :", t_amp)

    if device.type == "cuda":
        if t_amp < t_fp32:
            print("AMP is faster on GPU")
        else:
            print("FP32 is faster (depends on workload)")
```

```
...
--- Performance Summary ---
FP32 Time: 0.4018263816833496
AMP Time : 0.23004746437072754
AMP is faster on GPU
```

Variables Terminal 4:14 AM T4 (Python 3)

Fig 49:Performance Summary Output

Next, mixed precision using autodiff was tested to increase speed.

```
from torch.cuda.amp import autocast
```

```
with autocast():  
    result = A @ B
```

```
print("AMP computation completed")
```

Finally, quantization-aware training with PyTorch library was tested.

```
import torch.quantization as tq
```

```
model = torch.nn.Linear(10, 1)  
model.qconfig = tq.get_default_qat_qconfig("fbgemm")
```

```
tq.prepare_qat(model, inplace=True)  
tq.convert(model, inplace=True)
```

```
print("Quantization completed")
```

After this, quantization aware training was demonstrated using PyTorch libraries..

```
import torch.quantization as tq
```

```
model_q = torch.nn.Sequential(  
    torch.nn.Linear(10, 10),  
    torch.nn.ReLU(),  
    torch.nn.Linear(10, 1)  
)
```

```
model_q.qconfig = tq.get_default_qat_qconfig("fbgemm")  
tq.prepare_qat(model_q, inplace=True)
```

```
for _ in range(5):  
    data = torch.randn(4, 10)  
    out = model_q(data)
```

```
model_q.eval()  
tq.convert(model_q, inplace=True)
```

```
print("Quantized model ready")
```

There was an improvement in efficiency and reduced computation costs.

Task 9 : Accelerating Neural Network Inferencing: TensorRT & Triton Inference Server

Date: 28-04-2026

Procedure

This experiment was conducted to accelerate neural network inferencing using TensorRT and Triton Inference Server. It involved model conversion, repository creation, inference server deployment, and finally model inferencing.

Step-by-Step Explanation

A model was generated using PyTorch and exported to the ONNX format.

```
import os
os.environ["TORCH_ONNX_USE_DYNAMO"] = "0"

import torch
import torchvision.models as models

model = models.resnet18(pretrained=True)
model.eval()

dummy_input = torch.randn(1,3,224,224)

torch.onnx.export(model, dummy_input, "model.onnx", opset_version=11)
```

This model was then converted to the tensorRT engine.

```
trtexec --onnx=model.onnx --saveEngine=model.plan
```

The next step involved the generation of the model repository and placement of the engine files in the repository.

```
model_repository/
├── my_model/
│   └── 1/
│       └── model.plan
```

Configuration for this repository was made where the input/output information were specified.

```
name: "my_model"
platform: "tensorrt_plan"
max_batch_size: 1
```




Fig 50:Triton Docker Pull




Fig 51:ONNX Model Script




Fig 52:ONNX Model Created




Fig 53: Triton Config File




Fig 54:Triton Server Running and status

Finally, the Triton inference server was deployed using the Docker application.

```
docker run --gpus all -p8000:8000 -p8001:8001 -p8002:8002 -v C:\Users\afeef\model_repository:/models  
nvcr.io/nvidia/tritonserver:24.03-py3 tritonserver --model-repository=/models
```

Server status checked.

```
curl http://localhost:8000/v2/health/ready -UseBasicParsing
```

Model is successfully deployed and Inference speed is enhanced using GPU.

Task 10 : Monitoring Load Balancers & Schedulers using Docker with Jupyter, Prometheus, and Grafana

Date: 01-5-2026

Procedure

In this experiment, load balancing and scheduling services were deployed using Docker and performance metrics were analyzed using Prometheus and Grafana visualization tools. Computational load was applied using a Jupyter notebook.

Steps Involved

In this experiment, load balancing and scheduling services were deployed using Docker Compose and visualizations were carried out using Prometheus and Grafana visualization tooling. A Jupyter notebook was used for generating computational load.

Step-by-Step Explanation

Project directory was made, and docker-compose.yml was created.

```
mkdir exp10
cd exp10
```

Docker Compose configuration consisted of following components such as Jupyter, Prometheus, Grafana and Node Exporter.

```
version: '3'
services:
  jupyter:
    image: jupyter/base-notebook
    ports: ["8888:8888"]
```

```
prometheus:
  image: prom/prometheus
  ports: ["9090:9090"]
```

```
grafana:
  image: grafana/grafana
  ports: ["3000:3000"]
```

```
node-exporter:
  image: prom/node-exporter
  ports: ["9100:9100"]
```




Fig 60: Docker Compose Config




Fig 61:Prometheus Config File




Fig 63: Prometheus Target Health Status




Fig 64:Grafana Monitoring Dashboard




Fig 65:Grafana System Metrics Dashboard




Fig 65:Grafana System Metrics Dashboard




Fig 66:Jupyter Load Test Script

Services were started by using command:

```
docker compose up
```

The services were accessed through:

```
http://localhost:8888
```

```
http://localhost:9090
```

```
http://localhost:3000
```

And services can be accessed using following URLs.

```
import numpy as np
```

```
for i in range(200):
```

```
    a = np.random.rand(1500,1500)
```

```
    b = np.dot(a,a)
```

CPU usage was monitored in Grafana dashboard.

Product Innovation

Title: COVID-19 Detection System: AI-Powered Clinical Diagnostic Platform for COVID-19 Detection using Deep Learning and Explainable AI

1. Introduction

The presented product introduces an innovative AI-driven solution for the automated diagnosis of COVID-19 and related pulmonology conditions via chest X-ray images. Unlike other machine learning solutions that rely on one predictive algorithm, this system integrates several deep learning models, explainable AI components, and a number of scalable deployment options to ensure maximum interpretability, usability, and adaptability. In addition to delivering accurate predictions, the system's aim is to be reliable and useful to medical specialists when used in their clinical practices.

To train the proposed deep learning model, thousands of chest X-rays with labels were employed in this project. This data helps the model to identify several classes: COVID-19, normal lungs, viral pneumonia, and lung opacity. What makes this innovation different from traditional solutions is that this platform utilizes the power of a multi-model ensemble mechanism that involves the analysis of the same input image by several different models that reach a conclusion.

Along with improving the performance of the predictive process, this feature adds another layer of reliability to the diagnostic process since more than one model has to make the same decision for it to happen. Besides achieving the highest level of prediction accuracy, the system emphasizes the importance of being transparent and explainable. Clinicians expect a detailed explanation for each predicted class to help diagnose the patient and take further action.

A developed implementation of the current work is present at the following link:

<https://github.com/ssdafeef/PulmonaryAI>

2. Problem Statement

Current AI-based diagnostic solutions have faced some critical issues within the healthcare domain that limit their usage within clinical practice. One of the most prominent problems lies in using a single deep learning model that makes them prone to bias. Considering that any errors that the system makes will directly affect the health and life of patients, it is crucial to develop a reliable solution to deal with such a problem.

In addition, some diagnostic systems are primarily designed for research purposes and thus do not reflect real-life processes of working with medical data and making diagnoses. For example, they lack scalability and do not offer an opportunity to integrate the tool into existing systems and software solutions. In reality, there is a variety of steps required to be taken for diagnosing, such as data validation, collaboration, and structuring.

Another challenge of the current AI diagnostic solution is its inability to manage uncertainty. Current technologies tend to produce overconfident predictions that may influence practitioners'

decisions in a wrong way. Moreover, there is a lack of some practical features, like generating reports, classification of the disease severity level, and multi-language support.

Moreover, such systems do not use any approaches for uncertainty management or providing confidence scores, making predictions highly overconfident. Another important point is the lack of other useful features like automated report generation, severity categorization, or multilinguality, which makes these solutions not usable in practice. Hence, the aim of this project is to develop a system that will have all those shortcomings addressed.

3. Proposed Solution

The proposed system represents a hybrid AI solution incorporating the latest advancements in deep learning, explainable artificial intelligence, and deployment frameworks in order to deliver a full-fledged clinical decision support system for chest X-ray image assessment. As the core component of the system, a convolutional neural network with a special attention module is used. The chosen network structure is ResNet50.

For increased reliability of the solution, a multi-model ensemble method will be utilized. It will consist of several independent neural networks performing predictions based on a provided image. Then, using a certain algorithm, a consensus between the predictions generated by individual networks will be achieved. Thus, only the most confident predictions from different networks will make their way to the final result. This is inspired by the process of collaborative work of physicians in reality, where a team of specialists contributes to a final decision.

The problem of interpretability and explanation of predictions will also be solved using explainable artificial intelligence techniques. Grad-CAM algorithms will allow generating visual explanations represented by the heatmaps showing what parts of the image influenced the decision of the classifier. Moreover, a separate language model will be able to generate a natural-language explanation of a predicted diagnosis. It will be accompanied by the reasoning behind the prediction, its consequences, and possible further actions.

Automated generation of reports with all predictions and explanations will be implemented as another crucial feature. Reports generated by the solution will be in the PDF format, containing all data regarding the analysis of a chest X-ray image, including severity level. Moreover, there will be multilingual capabilities of the solution, along with a severity classification system.

4. System Workflow

The operation of the system involves a pipeline workflow simulation that resembles a realistic diagnostic scenario. The procedure starts by uploading a chest X-ray image through the web application. It is designed to allow easy interaction for both experienced and non-experienced users.

Once the image is uploaded, it is preprocessed in the backend environment. It includes image size rescaling, normalization of the image's pixel values, and conversion to a specific format. Proper preprocessing is vital for achieving consistent and reliable performance of AI models.

Further, the image is subjected to the analysis of several AI models. Each of them independently provides probability outputs related to different classes. Outputs of different models are combined using the technique of voting. If there is consistent behavior between the models, a certain classification result with high confidence level is obtained. However, in case of discrepancy, the

system informs about the inconsistency and suggests additional analysis. After the prediction, the system uses explainable AI and provides heatmaps highlighting the most influential areas in the image. Further, an AI-driven natural language explanation of the results is generated by the language model.

Additionally, the system performs severity classification for COVID-19 infection cases. As a final step, all information is aggregated and provided via the interface. A PDF document containing information related to the diagnosis is prepared automatically and ready for download.

5. Key Innovations

- AttentionResNet50 with a custom attention mechanism for COVID-19 classification
- Hybrid inference approach using NVIDIA Triton Inference Server and local Keras models
- Explainable AI based on the Grad-CAM visualization of every prediction
- Voting between multiple inference branches based on the confidence in prediction
- Ai-powered explanation using a language model (LLM; Google Gemini)
- Automated PDF report generation with multi-language support
- Four-class classification of Chest X-rays: COVID-19, Normal, Viral Pneumonia, Opacity
- Accuracy ~94%, ~0.97 AUC-ROC on 21,165 chest X-ray images
- Severity classification for COVID-19 infections (mild, moderate, severe)
- Fully deployed system using Docker Compose
- Model serving on NVIDIA Triton Inference Server
- End-to-end AI pipeline for uploading images and generating PDF reports




Figure 1: Architecture Diagram of model

Medical AI Diagnostic Portal

Agentic AI integration for hospital resource optimization and triage support.

Analyses
3

Status
Ready

Size
0.04 MB



Change Image

Run AI Prediction

Copy Summary

Clear Session

Diagnosis Results

High confidence

Classification:

COVID

AI Confidence:

91.98%

What To Do Next

Likely Condition: COVID-19 pattern suspected

Urgency: Prompt physician review recommended within the same day.

Recommended Actions

- Isolate patient as per local infection prevention protocol.
- Check oxygen saturation, respiratory rate, and temperature at triage.
- Correlate with symptoms and confirm using RT-PCR/rapid antigen as indicated.

Red Flags

- SpO2 below 94%
- Increasing shortness of breath
- Persistent high fever or chest pain

Follow-Up: Reassess clinically in 12 to 24 hours or earlier if symptoms worsen.

Figure 2: X-ray prediction output showcasing result


Medical AI Diagnostic Portal	123 Trichy, India
Doctor Name	Specialization
Dr. Olivia Greene	Doctor
Visit Date	Patient Name
05 / 03 / 2026	Sarah Anderson
Birth Date	Medical Number
01 / 01 / 1989	MA567891
IHI	Phone
5556-9669-9654-7788	+1 (555) 789-0123
Email	Hospital Contact
s.anderson@mail.com	+91 (XXX) 123-4567
Hospital Email	Hospital Site
info@Medical_AI_Diagnostic.com	www.Medical_AI_Diagnostic.com

Source: remote-llm | Latency: 9296 ms | Gemini-backed LLM

Impression: Multifocal, predominantly peripheral ground-glass opacities and consolidations are present bilaterally, significantly more pronounced in the left lung. These findings are highly suggestive of viral pneumonia, such as COVID-19, given the clinical context and high AI confidence. Correlation with RT-PCR testing and clinical symptoms is strongly recommended.

Narrative Summary: The frontal chest radiograph demonstrates extensive, patchy opacities distributed throughout both lung fields. There is a notable peripheral and mid-to-lower zone predominance, which is characteristic of atypical or viral pneumonias. The left lung shows more confluent areas of consolidation compared to the right, particularly in the mid-lung zone. No significant pleural effusions or pneumothorax are identified. The cardiomedial silhouette remains within normal limits for size and configuration, and the visualized osseous structures are unremarkable.

LLM Pattern Map



- Ground-glass opacity: Patchy peripheral opacities in the right mid-to-

Figure 3: The LLM model showcasing overlay of detected area




Figure 4: The Whole Project running on docker

PERSONAL

Search

Ctrl+K

A

Gordon BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit BETA

Docker Hub

Docker Scout

Extensions

Upgrade plan

Engine running

RAM 3.46 GB CPU 0.37% Disk: 66.77 GB used (limit 1006.85 GB)

Terminal v4.71.0

seaaaai

D:\seaaaai

View configurations

api	INFO:	127.0.0.1:40050	- "GET /health HTTP/1.1" 200 OK
afeefahmed/seaaaai:8010:8010	api-client	INFO:	127.0.0.1:52096 - "GET /health HTTP/1.1" 200 OK
		INFO:	127.0.0.1:53560 - "GET /health HTTP/1.1" 200 OK
triton-server	api	INFO:	127.0.0.1:40076 - "GET /health HTTP/1.1" 200 OK
seaaaai-triton-server:8000:8000	api-client	INFO:	127.0.0.1:56450 - "GET /health HTTP/1.1" 200 OK
		INFO:	127.0.0.1:50136 - "GET /health HTTP/1.1" 200 OK
api-client	api-client	INFO:	127.0.0.1:49298 - "GET /health HTTP/1.1" 200 OK
seaaaai-api-client:8011:8001	api	INFO:	127.0.0.1:53052 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:44514 - "GET /health HTTP/1.1" 200 OK
		INFO:	127.0.0.1:43372 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:43338 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:36108 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:49668 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:37116 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:34674 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:59560 - "GET /health HTTP/1.1" 200 OK
		INFO:	127.0.0.1:60654 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:48212 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:40018 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:59322 - "GET /health HTTP/1.1" 200 OK

Figure 5: All the container inside docker container

1. Dataset and Training Details

The creation of the proposed COVID-19 Detection System is possible because of the [COVID-19 Radiography Database](#), which is the commonly-used dataset for medical image classification tasks. It includes numerous chest radiography images grouped into several classes: COVID-19, Normal, Viral Pneumonia, and Lung Opacity. The presence of over 21,000 images makes it diverse enough to train the system with deep learning methods.

The dataset has training, validation, and testing samples to ensure the possibility of reliable evaluation of the system's performance. For example, the training subset is utilized to train models and learn necessary patterns, while the validation one monitors the learning process and prevents overfitting. At last, test data is used to assess the results and estimate the performance in real-world conditions.

To enhance the performance of models, multiple preprocessing techniques are applied. They include resizing images into standard input dimensions (224×224 pixels) and pixel values normalization. Some other techniques used are rotations, flips, and scaling of images to diversify data. Finally, because of class imbalance, class weighting was applied during the learning stage to ensure fair results.

Deep learning models are trained using TensorFlow library in Google Colab equipped with GPU. Transfer learning is used to pretrain the models on pretrained architectures such as ResNet50 and DenseNet121. The best performance of models trained with transfer learning and without it are compared. Moreover, class weighting and standard training techniques are also explored. The accuracy score reaches about 94% when using the most productive approach. Thus, the selected combination of techniques yields impressive results and ensures an accurate system.

In general, with an extensive dataset, numerous preprocessing techniques, transfer learning, and GPU utilization, the developed system is highly accurate and reliable.

2. Deployment and Scalability

To provide scalability and portability to the developed system, a number of techniques can be used. They include the creation of self-contained units, packaging of components into containers, and implementing cloud computing capabilities. All mentioned elements are present in the design of the system, making it scalable and portable.

A high-performance inference server is utilized to efficiently process numerous requests and perform in real-time environment. Additionally, because of the hybrid architecture, the system works in environments with or without GPUs. The use of container orchestration makes it easy to scale the system up to handle higher loads.

By implementing such technologies as container orchestration, the system can easily scale to process higher volumes of data. Therefore, the solution can be deployed in hospitals, scientific centers, and even in remote medical centers. Generally speaking, the deployment approach enables the creation of an AI tool that is powerful and pragmatic.

3. Real-World Impact

With its help, the diagnosis of such diseases as coronavirus infection, lung inflammation, and other related conditions can be carried out with precision and accuracy. Furthermore, the AI-driven diagnostic tool will make it easier for specialists to make informed choices since the system offers both the results and explanations. It may serve as an additional aid to those parts of the world where experienced radiologists cannot be found.

The use of such a system can positively affect healthcare through improved speed and quality of diagnostics. Moreover, the possibility of generating structured reports and evaluating the severity of symptoms increases the efficiency of the entire diagnostic process.

4. Conclusion

This paper has offered an insight into developing a powerful system designed to help people get reliable diagnoses. Such a tool combines various modern approaches used in machine learning and can provide useful solutions for medical diagnostics.

The project can be seen as an example of the successful application of artificial intelligence in solving complex problems and creating helpful applications for everyday practice.

Review 1: Quality of the Module Developed

First of all, an essential aspect of the developed module is the presence of a structure in the system and the successful use of several advanced technologies within the model. It should be mentioned that the process of creating the model was performed on the basis of a rather rich COVID-19 Radiography Database, which includes many samples of chest radiograph images. The interaction with such a resource allowed considering important aspects, such as class distributions, data preparation, and training processes, contributing greatly to the creation of the proposed system.

Another crucial feature of the module is high computational efficiency due to the possibility to work with GPU. Training models in Google Colab using GPU resources allowed performing faster tests and avoiding problems related to the necessity to process big amounts of data and train complicated structures of deep learning models.

Moreover, it is important to note that another aspect of the quality of the module concerns its architecture. In the context of the described project, one can observe how the system successfully combines such important components as data preprocessing, model prediction, explaining AI, and report generation. Moreover, each component performs its function while improving the performance of the overall system.

In addition to the previous aspects, it is vital to note the presence of the feature of the generated explanations by the system. The ability to see visual explanations of results and use the information provided by AI about the disease increases the quality of the tool significantly.

Overall, it should be noted that the created project demonstrates the high quality of its modules since theoretical ideas are combined with the implementation of the solution..

Create

Home

Competitions

Benchmarks

Game Arena

Data Hub

More

Your Work

Viewed

COVID-19 Radiograp...

Chest X-Ray Images ...

finalupdated

tbcxr

Pneumonia Detection...

Edited

finalupdated

View Active Events

Search

COVID-19 Radiography Database

1194

Code

Download

Data Card

Code (511)

Discussion (27)

Suggestions (0)

*COVID-19 CHEST X-RAY DATABASE

A team of researchers from Qatar University, Doha, Qatar, and the University of Dhaka, Bangladesh along with their collaborators from Pakistan and Malaysia in collaboration with medical doctors have created a database of chest X-ray images for COVID-19 positive cases along with Normal and Viral Pneumonia images. This COVID-19, normal and other lung infection dataset is released in stages. In the first release we have released 219 COVID-19, 1341 normal and 1345 viral pneumonia chest X-ray (CXR) images. In the first update, we have increased the COVID-19 class to 1200 CXR images. In the 2nd update, we have increased the database to 3616 COVID-19 positive cases along with 10,192 Normal, 6012 Lung Opacity (Non-COVID lung infection) and 1345 Viral Pneumonia images. We will continue to update this database as soon as we have new x-ray images for COVID-19 pneumonia patients.

**COVID-19 data:

COVID data are collected from different publicly accessible dataset, online sources and published papers.

-2473 CXR images are collected from padchest dataset[1].

-183 CXR images from a Germany medical school[2].

COVID

2 directories

Lung_Opacity

2 directories

Normal

2 directories

Viral Pneumonia

2 directories

README.md.txt

3.35 kB

COVID.metadata.xlsx

Lung_Opacity.metadat...

Normal.metadata.xlsx

Viral Pneumonia.meta...

COVID-1013.png

COVID-1016.png

COVID-1017.png

COVID-1018.png

COVID-1019.png

COVID-102.png

COVID-1020.png

COVID-1021.png

COVID-1022.png

COVID-1023.png

COVID-1024.png

3586 more

Lung_Opacity

Normal

Viral Pneumonia

README.md.txt

COVID.metadata.xlsx

Lung_Opacity.metadat

Normal.metadata.xlsx

Viral Pneumonia.metac

Summary

42.3k files

16 columns

Figure 6: All the different dataset images




Figure 7: Untrained model false detection before fine tuning

Step 1. Original X-ray



Step 2. Grad-CAM novelty



Figure 8: Huge detection improvement after fine-tuning the model

Review 2: Containerization and Deployment

One of the most important features of the developed COVID-19 Detection System is the presence of clear mechanisms for containerization and deployment which make the system easily reproducible, portable and fit for productive use.

During the process of system creation, a special focus was put on following modern software architecture and deployment trends. Therefore, containerization was used in accordance with the modern principles based on Docker and Docker Compose technologies.

System Architecture:

The developed COVID-19 Detection System can be viewed as a multi-component architecture including the following parts:

- **Frontend:** React app that uses Vite for development and optimized building
- **Backend:** FastAPI server for images preprocessing, orchestration and API management
- **Model Inference:** NVIDIA Triton Inference Server for efficient, high-performance inference of models
- **Local Model:** As a fallback, Keras model can be used to ensure more reliable results

All of these components were packaged separately into containers by means of Docker technology. It guarantees that no external dependencies exist, and all necessary libraries are packed.

Docker Compose Orchestration

In order to facilitate the launch of multiple services, Docker Compose technology was used. The file `docker-compose.yml` describes services, their configuration, dependencies and networking. The system can be started with the following command:

```
docker compose up -d
```

This command does the following:

- Building and launching of Triton Inference Server with COVID classifier model loaded
- Starting a FastAPI server with correct health check setup and dependencies
- Launching React application
- Setting up inter-service communication and volumes for model loading

NVIDIA Triton Model Serving

When deploying a system, production-ready NVIDIA Triton Inference Server is used for efficient inference of pre-trained models. For setting up the service, the corresponding Docker image from NVIDIA repository is downloaded first:

```
docker pull nvcr.io/nvidia/tritonserver:23.10-py3
```

Service Endpoints

After deployment, all service endpoints can be accessed using the following URLs:

- **Frontend:** <http://localhost:3000>— React Vite application for uploading X-rays and visualizing results
- **Backend API:** <http://localhost:8010/docs> — FastAPI documentation with prediction endpoint
- **Triton Server:** <http://localhost:8001>— NVIDIA Triton Inference Server for models inference (HTTP), <http://localhost:8002> (gRPC)

Benefits of Deployment

The following are the key benefits of this deployment solution:

- **Consistency:** Standardized execution in development, testing, and production stages
- **Mobility:** All services can be executed on any device that supports Docker and Docker Compose
- **Scaleability:** Individual scaling of services based on demand (front-end, back-end, inference)
- **Resilience:** Health checks guarantee the availability of services; fault tolerance is provided by the mixed inference pipeline using Triton and local Keras
- **Deployment Ready:** Enterprise-level orchestration that is capable of being deployed to cloud-based systems or on-site infrastructure

Medical AI Diagnostic Portal

Agentic AI integration for hospital resource optimization and triage support.

Analyses
3

Status
Ready

Size
0.04 MB



Change Image

Run AI Prediction

Copy Summary

Clear Session

Diagnosis Results

High confidence

Classification:

COVID

AI Confidence:

91.98%

What To Do Next

Likely Condition: COVID-19 pattern suspected

Urgency: Prompt physician review recommended within the same day.

Recommended Actions

- Isolate patient as per local infection prevention protocol.
- Check oxygen saturation, respiratory rate, and temperature at triage.
- Correlate with symptoms and confirm using RT-PCR/rapid antigen as indicated.

Red Flags

- SpO2 below 94%
- Increasing shortness of breath
- Persistent high fever or chest pain

Follow-Up: Reassess clinically in 12 to 24 hours or earlier if symptoms worsen.

Figure 9: Main Dashboard


Medical AI Diagnostic Portal	123 Trichy, India
Doctor Name	Specialization
Dr. Olivia Greene	Doctor
Visit Date	Patient Name
05/03/2026	Sarah Anderson
Birth Date	Medical Number
01/01/1989	MA567891
IHI	Phone
5556-9669-9654-7788	+1 (555) 789-0123
Email	Hospital Contact
s.anderson@mail.com	+91 (XXX) 123-4567
Hospital Email	Hospital Site
info@Medical_AI_Diagnostic.com	www.Medical_AI_Diagnostic.com

Source: remote-llm | Latency: 9296 ms | Gemini-backed LLM

Impression: Multifocal, predominantly peripheral ground-glass opacities and consolidations are present bilaterally, significantly more pronounced in the left lung. These findings are highly suggestive of viral pneumonia, such as COVID-19, given the clinical context and high AI confidence. Correlation with RT-PCR testing and clinical symptoms is strongly recommended.

Narrative Summary: The frontal chest radiograph demonstrates extensive, patchy opacities distributed throughout both lung fields. There is a notable peripheral and mid-to-lower zone predominance, which is characteristic of atypical or viral pneumonias. The left lung shows more confluent areas of consolidation compared to the right, particularly in the mid-lung zone. No significant pleural effusions or pneumothorax are identified. The cardiomedial silhouette remains within normal limits for size and configuration, and the visualized osseous structures are unremarkable.

LLM Pattern Map



- **Ground-glass opacity:** Patchy peripheral opacities in the right mid-to-lower lung zone.

Figure 10: LLM model Insights

LLM Annotated Image




Figure 11: Model Overlay of Detected virus

:

PERSONAL

Search

Ctrl+K

?

4

A

-

×

Gordon BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit BETA

Docker Hub

Docker Scout

Extensions

Upgrade plan

<

seaaai

D:\seaaai

View configurations

api	INFO:	127.0.0.1:40050	- "GET /health HTTP/1.1" 200 OK
api afeefahmed/seaaai 8010:8010	api-client	INFO:	127.0.0.1:52096 - "GET /health HTTP/1.1" 200 OK
		INFO:	127.0.0.1:53560 - "GET /health HTTP/1.1" 200 OK
triton-server	INFO:	127.0.0.1:40076	- "GET /health HTTP/1.1" 200 OK
seaaai-triton-server 8000:8000 Show all ports (3)	api-client	INFO:	127.0.0.1:56450 - "GET /health HTTP/1.1" 200 OK
api-client	INFO:	127.0.0.1:49298	- "GET /health HTTP/1.1" 200 OK
seaaai-api-client 8011:8001	api	INFO:	127.0.0.1:53052 - "GET /health HTTP/1.1" 200 OK
frontend	INFO:	127.0.0.1:44514	- "GET /health HTTP/1.1" 200 OK
seaaai-frontend 3000:80		INFO:	127.0.0.1:43372 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:43338 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:36108 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:49668 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:37116 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:34674 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:59560 - "GET /health HTTP/1.1" 200 OK
		INFO:	127.0.0.1:60654 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:48212 - "GET /health HTTP/1.1" 200 OK
	api-client	INFO:	127.0.0.1:40018 - "GET /health HTTP/1.1" 200 OK
	api	INFO:	127.0.0.1:59322 - "GET /health HTTP/1.1" 200 OK

Engine running

RAM 3.46 GB CPU 0.37% Disk: 66.77 GB used (limit 1006.85 GB)

Terminal v4.71.0

Figure 12: All containers running in docker